

Vansh Sharma
1NT19IS181
C2 Batch

Date: 23/06/2022

Experiment no. 5

Download the file as .csv file. Using the Sales.csv file, create a database called “salesdb”. Create a table under that with the heading “sales”. (Adjust the rows by looking at the csv file columns)

1. Insert 5 records using Insert command in HiveQL

```
hive> insert into cred values("Product1",1500,"Axis","ABC","Bangalore","karnataka","India");
Query ID = hdoop_20220622150101_c4303631-0575-461e-a423-662c053b22a1
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1655888206664_0001, Tracking URL = http://admin1-HF-280-04-HF-Business-PC:8088/proxy/application_1655888206664_0001/
```

```
2022-06-22 15:01:20,218 Stage-1 Map = 100%, reduce = 100%, cumulative CPU 3.97 sec
MapReduce Total cumulative CPU time: 3 seconds 970 msec
Ended Job = job_1655888206664_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/cred/.hive-staging_hive_2022-06-22_15-01-01_526_23007417153159
37572-1/-ext-10300
Loading data to table default.cred
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduces: 1 Cumulative CPU: 3.97 sec HDFS Reads: 21415 HDFS Writes: 457 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 970 msec
OK
Time taken: 20.241 seconds
```

```
hive> select * from cred;
OK
Product1      1500    Axis    ABC    Bangalore    karnataka    India
Time taken: 0.15 seconds, Fetched: 1 row(s)
hive> █
```

2. Import the dataset from the sales.csv file

```
hive> load data local inpath '/home/hdoop/Downloads/sales_withoutHeader - sales_withoutHeader.csv' into table cred;
Loading data to table default.cred
OK
Time taken: 0.179 seconds
hive> █
```

```
hive> select * from cred;
OK
Product1      1500   Axis   ABC   Bangalore   karnataka   India
Product NULL  Payment_Type Name   City   State   Country
Product1      1200   Mastercard   carolina   Basildon   England United Kingdom
Product1      1200   Visa   Betina   Parkville   MO   United States
Product1      1200   Mastercard   Federica e Andrea   Astoria OR   United States
Product1      1200   Visa   Gouya   Echuca   Victoria   Australia
Product2      3600   Visa   Gerd W   Cahaba Heights   AL   United States
Product1      1200   Visa   LAURENCE   Mickleton   NJ   United States
Product1      1200   Mastercard   Fleur   Peoria   IL   United States
Product1      1200   Mastercard   adam   Martin   TN   United States
Product1      1200   Mastercard   Renee Elisabeth   Tel Aviv   Tel Aviv   Israel
Product1      1200   Visa   Aidan   Chatou   Ile-de-France   France
Product1      1200   Diners   Stacy   New York   NY   United States
Product1      1200   Amex   Heidi   Eindhoven   Noord-Brabant   Netherlands
Product1      1200   Mastercard   Sean   Shavano Park   TX   United States
Product1      1200   Visa   Georgia Eagle   ID   United States
Product1      1200   Visa   Richard   Riverside   NJ   United States
Product1      1200   Diners   Leanne   Julianstown   Meath   Ireland
Product1      1500   Visa   Janet   Ottawa   Ontario   Canada
```

3. Count the number of sales done by each country

```
hive> select country,count(*) from cred group by country;
Query ID = hdoop_20220622150858_34785f0a-cb37-4cf0-95e2-34bfe5b8b97f
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1655088206664_0003, Tracking URL = https://admin1-HP-280-04-MT-Business-PC:8088/proxy/application_1655088206664_0003/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1655088206664_0003
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-06-22 15:09:03,186 Stage-1 map = 0%, reduce = 0%
2022-06-22 15:09:09,281 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.32 sec
2022-06-22 15:09:13,429 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.75 sec
MapReduce Total cumulative CPU time: 2 seconds 750 msec
Ended Job = job_1655088206664_0003
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.75 sec HDFS Read: 70928 HDFS Write: 1444 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 750 msec
OK
```

Japan	2	Argentina	1
Jersey	1	Australia	38
Kuwait	1	Austria	7
Latvia	1	Bahrain	1
Luxembourg	1	Belgium	8
Malaysia	1	Bermuda	1
Malta	2	Brazil	5
Mauritius	1	Bulgaria	1
Moldova	1	CO	1
Monaco	2	Canada	76
Netherlands	22	Cayman Isla	1
New Zealand	6	China	1
Norway	16	Costa Rica	1
Philippines	2	Country	1
Poland	2	Czech Republic	3
Romania	1	Denmark	15
Russia	1	Dominican Republic	1
South Africa	5	Finland	2
South Korea	1	France	27
Spain	12	Germany	25
Sweden	13	Greece	1
Switzerland	36	Guatemala	1
Thailand	2	Hong Kong	1
The Bahamas	2	Hungary	3
Turkey	6	Iceland	1
Ukraine	1	India	3
United Arab Emirates	4		
United Kingdom	108		
United States	442		

Time taken: 16.707 seconds, Fetched: 58 row(s)
hive>

4. Count the number of sales done by each state

```

hive> select state,count(*) from cred group by state;
Query ID = hdoop_20220622151306_b74a958b-c33b-4aa4-9a2c-c3936e728b1
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reducers=<number>
Starting Job = job_1655888206664_0004, Tracking URL = http://admini-wr-260-c4-wr-business-rc:8088/proxy/application\_1655888206664\_0004/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1655888206664_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-06-22 15:13:05,231 Stage-1 map = 0%, reduce = 0%
2022-06-22 15:13:09,395 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.32 sec
2022-06-22 15:13:13,527 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.77 sec
MapReduce Total cumulative CPU time: 2 seconds 770 msec
Ended Job = job_1655888206664_0004
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 cumulative CPU: 2.77 sec HDFS Read: 71010 HDFS Write: 4814 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 770 msec
OK

```

Bern	1	Georgia	1
Black River	1	Groningen	2
Bohol	1	Guangdong	1
British Columbia	18	Guatemala	1
Brittany	2	HI	9
Brussels (Bruxelles)	6	Hainaut	1
Bucuresti	1	Hamburg	1
Budapest	1	Hamilton	1
Buenos Aires	1	Haryana	1
Buskerud	1	Heredia	1
CA	66	Hessen	2
CO	10	Hordaland	2
CT	9	IA	3
Catalonia	3	ID	2
Ceara	1	IL	16
Centennial	1	IN	2
Central Bohemia	1	Ile-de-France	8
Chisinau	1	Istanbul	5
Cork	10	Ita-Suomen Laani	1
DC	3	Izmir	1
DE	1	KS	2
Distrito Nacional	1	NY	2
Dubayy	5	Nerry	1
Dublin	22	Kiev	1
East Bohemia	1	Kilkenny	1
England	86	Kobenhavn	7
Etelä-Suomen Laani	1	Krung Thep	1
FL	29	Kuala Lumpur	1
Franche-Comte	2	KwaZulu-Natal	2
Frederiksborg	3	LA	2
Fresport	1	Lazio	1
GA	21	Limburg	2

5. Display (Product, name) grouped by product

```

hive> select prod_id,count(uname) from cred group by prod_id;
Query ID = hdoop_20220622132000_6477cbcf-169c-439f-9dc3-bed12a0c428
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reducers=<number>
Starting Job = job_1655888206664_0005, Tracking URL = http://admini-SP-280-G4-MT-Business-PC:8088/proxy/application_1655888206664_0005/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1655888206664_0005
Hadoop Job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-06-22 15:20:06,711 Stage-1 map = 0%, reduce = 0%
2022-06-22 15:20:10,892 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.11 sec
2022-06-22 15:20:14,896 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.52 sec
MapReduce Total cumulative CPU time: 2 seconds 320 msec
Ended Job = job_1655888206664_0005
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.52 sec HDFS Read: 7126? HDFS Write: 183 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 320 msec
OK

```

```

OK
Product 1
Product1      848
Product2      136
Product3       15
Time taken: 16.848 seconds, Fetched: 4 row(s)

```

6. Create separate views for VISA and Mastercard

```

hive> create view visa as (
  > select * from sales where payment_type='Visa');
OK
Time taken: 1.437 seconds
hive> select * from visa;
OK
Product1      1200   Visa   Betina   Parkville   MD   United States
Product1      1200   Visa   Gouya   Echuca Victoria   Australia
Product2      3600   Visa   Gerd W   Cahaba Heights AL   United States
Product1      1200   Visa   LAURENCE   Mickleton   NJ   United States
Product1      1200   Visa   Aidan   Chatou   Ile-de-France   France
Product1      1200   Visa   Georgia   Eagle ID   United States
Product1      1200   Visa   Richard   Riverside   NJ   United States
Product1      1200   Visa   Janet   Ottawa Ontario Canada
Product2      3600   Visa   Sabine   London England United Kingdom
Product1      1200   Visa   Jeremy   Manchester England United Kingdom
Product1      1200   Visa   asuman   Chula Vista CA   United States
Product1      1200   Visa   Lisa   Sugar Land TX   United States
Product1      1200   Visa   chris   London England United Kingdom

hive> create view mastercard as (
  > select * from sales where payment_type='Mastercard');
OK
Time taken: 0.12 seconds

```

```
hive> select * from mastercard;
OK
Product1      1200    Mastercard    carolina      Basildon      England
United Kingdom
Product1      1200    Mastercard    Federica e Andrea    Astoria OR    U
nited States
Product1      1200    Mastercard    Fleur    Peoria    IL    United States
Product1      1200    Mastercard    adam    Martin    TN    United States
Product1      1200    Mastercard    Renee Elisabeth    Tel Aviv    Tel Avi
v    Israel
Product1      1200    Mastercard    Sean    Shavano Park    TX    United
States
Product1      1200    Mastercard    Nicola    Roodepoort    Gauteng South A
frica
Product1      1200    Mastercard    Lena    Kuopio    Ita-Suomen Laani    F
inland
Product1      1200    Mastercard    Katherine    New York    NY    U
nited States
Product1      1200    Mastercard    Linda    Miami    FL    United States
Product1      1200    Mastercard    SYLVIA    Vesenz    Geneve    Switzerland
Product1      1200    Mastercard    Stephanie    Badhoevedorp    Noord-H
olland Netherlands
Product1      1200    Mastercard    Anneli    Houston TX    United States
```

7. Show all the transactions done in Seattle

```
hive> select * from cred where city="Seattle";
OK
Product1      1200    Visa    Vicki and Bill    Seattle WA    United States
Product1      1200    Visa    Craig    Seattle WA    United States
Product1      1200    Visa    Kim    Seattle WA    United States
Product2      3600    Visa    cristina    Seattle WA    United States
Product1      1200    Amex    Justine    Seattle WA    United States
Product1      1200    Visa    Reto    Seattle WA    United States
Time taken: 0.103 seconds, Fetched: 6 row(s)
hive>
```

8. Find the max number of transactions done within the state of Ontario


```

hive> select payment_type,max(cnt) from (select payment_type,count(*) as cnt fr
om sales where state='Ontario' group by payment_type) as t;
FAILED: SemanticException [Error 10025]: Line 1:7 Expression not in GROUP BY ke
y 'payment_type'
hive> (select payment_type,count(*) as cnt from sales where state='Ontario' gro
up by payment_type order by cnt LIMIT 1);
Query ID = hdoop_20220622152136_ddb069b4-c015-452f-8e8d-d4c80854f2d2
Total jobs = 2
Launching Job 1 out of 2
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1655888018522_0011, Tracking URL = http://admin1-HP-280-G4-M
T-Business-PC:8088/proxy/application_1655888018522_0011/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1655888018522
_0011
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1

```

```

In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1655888018522_0014, Tracking URL = http://admin1-HP-280-G4-M
T-Business-PC:8088/proxy/application_1655888018522_0014/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1655888018522
_0014
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2022-06-22 15:22:51,659 Stage-2 map = 0%,  reduce = 0%
2022-06-22 15:22:55,789 Stage-2 map = 100%,  reduce = 0%, Cumulative CPU 1.06 s
ec
2022-06-22 15:22:59,900 Stage-2 map = 100%,  reduce = 100%, Cumulative CPU 2.34
sec
MapReduce Total cumulative CPU time: 2 seconds 340 msec
Ended Job = job_1655888018522_0014
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1  Reduce: 1   Cumulative CPU: 3.52 sec   HDFS Read: 71343
HDFS Write: 196 SUCCESS
Stage-Stage-2: Map: 1  Reduce: 1   Cumulative CPU: 2.34 sec   HDFS Read: 7790 H
DFS Write: 107 SUCCESS
Total MapReduce CPU Time Spent: 5 seconds 860 msec
OK
Visa      20
Time taken: 36.4 seconds, Fetched: 1 row(s)

```

9. Find the number of transactions whose price is in between 1500-3600

```

hive> select count(*) from sales where price>1500 and price<3600;
Query ID = hdoop_20220622152647_38f55533-4ad2-472b-97c2-a98792740b23
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1655888018522_0015, Tracking URL = http://admin1-HP-280-G4-M
T-Business-PC:8088/proxy/application_1655888018522_0015/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1655888018522
_0015
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-06-22 15:26:53,331 Stage-1 map = 0%, reduce = 0%

In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1655888018522_0015, Tracking URL = http://admin1-HP-280-G4-M
T-Business-PC:8088/proxy/application_1655888018522_0015/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1655888018522
_0015
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-06-22 15:26:53,331 Stage-1 map = 0%, reduce = 0%
2022-06-22 15:26:57,455 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.3 se
c
2022-06-22 15:27:03,601 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.02
sec
MapReduce Total cumulative CPU time: 4 seconds 20 msec
Ended Job = job_1655888018522_0015
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.02 sec HDFS Read: 72450
HDFS Write: 101 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 20 msec
OK
2
Time taken: 17.39 seconds, Fetched: 1 row(s)
hive>

```

10. List all the transactions done in the United States using Mastercard

```
hive> select * from sales where country="United States" and payment_type="Mastercard";
OK
Product1      1200      Mastercard      Federica e Andrea      Astoria OR      U
nited States
Product1      1200      Mastercard      Fleur      Peoria      IL      United States
Product1      1200      Mastercard      adam      Martin      TN      United States
Product1      1200      Mastercard      Sean      Shavano Park      TX      United
States
Product1      1200      Mastercard      Katherine      New York      NY      U
nited States
Product1      1200      Mastercard      Linda      Miami      FL      United States
Product1      1200      Mastercard      Anneli      Houston      TX      United States
Product1      1200      Mastercard      Marie      Ball Ground      GA      United
States
Product1      1200      Mastercard      Nikki      New Rochelle      NY      United
States
Product1      1200      Mastercard      June      Beachwood      OH      United
States
Product1      1200      Mastercard      chrissy W      Lebanon      NH      United
States
Product2      3600      Mastercard      Valda      Irvine      CA      United States
```